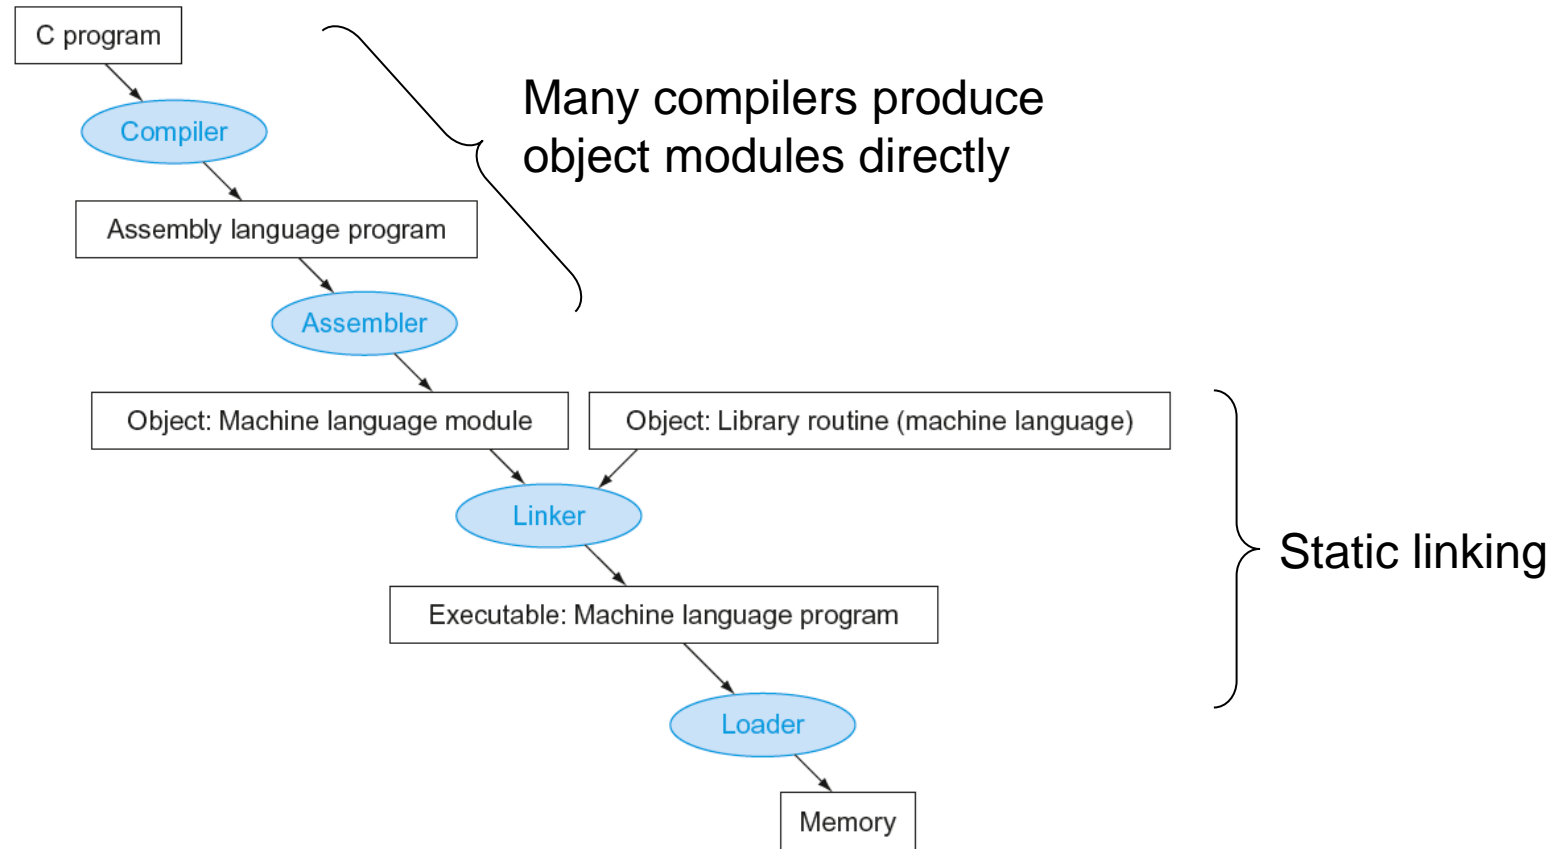


Fundamentals of Computer Architecture

ARM - Review

Teacher: Lobar Asretdinova

Translation and Startup



Producing an Object Module

- Assembler (or compiler) translates program into machine instructions
- Provides information for building a complete program from the pieces
 - Header: described contents of object module
 - Text segment: translated instructions
 - Static data segment: data allocated for the life of the program
 - Relocation info: for contents that depend on absolute location of loaded program
 - Symbol table: global definitions and external refs
 - Debug info: for associating with source code

Linking Object Modules

- Produces an executable image
 1. Merges segments
 2. Resolve labels (determine their addresses)
 3. Patch location-dependent and external refs
- Could leave location dependencies for fixing by a relocating loader
 - But with virtual memory, no need to do this
 - Program can be loaded into absolute location in virtual memory space

Loading a Program

- Load from image file on disk into memory
 1. Read header to determine segment sizes
 2. Create virtual address space
 3. Copy text and initialized data into memory
 - Or set page table entries so they can be faulted in
 4. Set up arguments on stack
 5. Initialize registers (including SP, FP)
 6. Jump to startup routine
 - Copies arguments to X0, ... and calls main
 - When main returns, do exit syscall

Dynamic Linking

- Only link/load library procedure when it is called
 - Requires procedure code to be relocatable
 - Avoids image bloat caused by static linking of all (transitively) referenced libraries
 - Automatically picks up new library versions

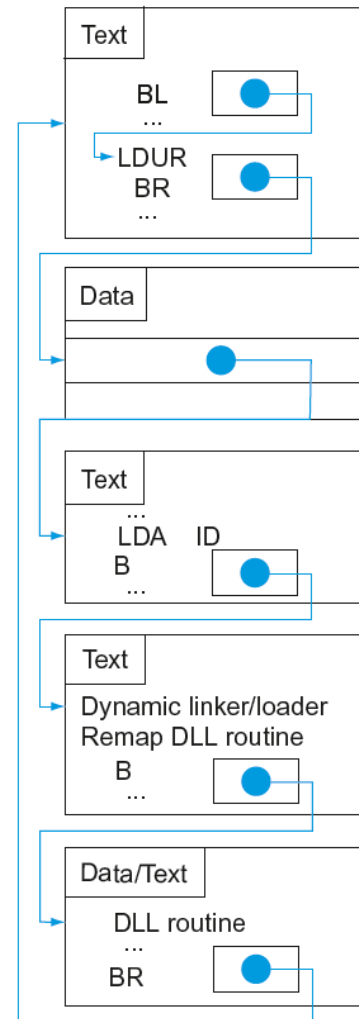
Lazy Linkage

Indirection table


Stub: Loads routine ID,
Jump to linker/loader

Linker/loader code

Dynamically
mapped code

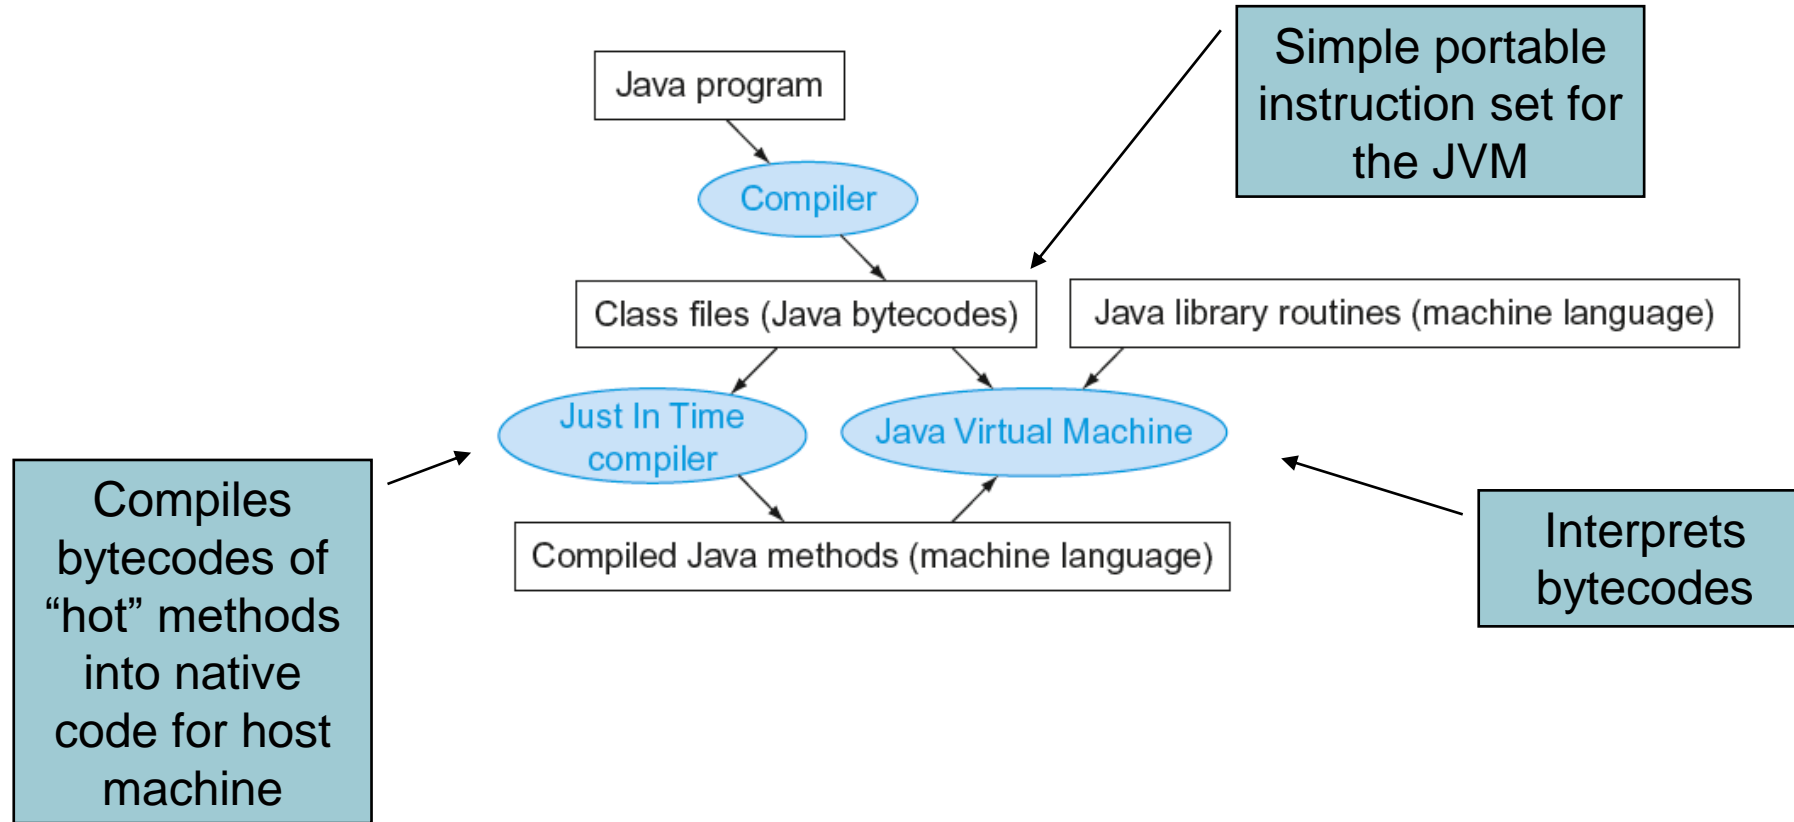


(a) First call to DLL routine



(b) Subsequent calls to DLL routine

Starting Java Applications



C Sort Example

- Illustrates use of assembly instructions for a C bubble sort function
- Swap procedure (leaf)

```
void swap(long long int v[],  
long long int k)  
{  
    long long int temp;  
    temp = v[k];  
    v[k] = v[k+1];  
    v[k+1] = temp;  
}
```

- v in X0, k in X1, temp in X9

The Procedure Swap

```
swap: LSL X10,X1,#3      // X10 = k * 8
      ADD X10,X0,X10     // X10 = address of v[k]
      LDR X9,[X10,#0]    // X9 = v[k]
      LDR X11,[X10,#8]   // X11 = v[k+1]
      STR X11,[X10,#0]   // v[k] = X11 (v[k+1])
      STR X9,[X10,#8]    // v[k+1] = X9 (v[k])
      BR LR              // return to calling routine
```

The Sort Procedure in C

- Non-leaf (calls swap)

```
void sort (long long int v[], size_t n)
{
    size_t i, j;
    for (i = 0; i < n; i += 1) {
        for (j = i - 1;
             j >= 0 && v[j] > v[j + 1];
             j -= 1) {
            swap(v, j);
        }
    }
}
```

- v in X0, n in X1, i in X19, j in X20

2, 5, 6, 3, 4
0 1, 2

n=5

1)

2) i=1

j=0

3) i=2

j=1

v[1]=5

v[2]=6

4) i=3 j=2

v[2]=6 v[3]=3

The Outer Loop

- Skeleton of outer loop:

- for (i = 0; i < n; i += 1) {

```
    MOV X19,XZR                // i = 0
for1tst:
    CMP X19, X1                // compare X19 to X1 (i to n)
    B.GE exit1                // go to exit1 if X19 ≥ X1 (i ≥ n)
```

(body of outer for-loop)

```
    ADD X19,X19,#1            // i += 1
    B for1tst                // branch to test of outer loop
exit1:
```

The Inner Loop

- Skeleton of inner loop:

- for (j = i - 1; j >= 0 && v[j] > v[j + 1]; j -= 1) {

```
    SUB X20, X19, #1
```

```
// j = i - 1
```

```
for2tst: CMP X20, XZR
```

```
// compare X20 to 0 (j to 0)
```

```
    B.LT exit2
```

```
// go to exit2 if X20 < 0 (j < 0)
```

```
    LSL X10, X20, #3
```

```
// reg X10 = j * 8
```

```
    ADD X11, X0, X10
```

```
// reg X11 = v + (j * 8)
```

```
    LDR X12, [X11, #0]
```

```
// reg X12 = v[j]
```

```
    LDR X13, [X11, #8]
```

```
// reg X13 = v[j + 1]
```

```
    CMP X12, X13
```

```
// compare X12 to X13
```

```
    B.LE exit2
```

```
// go to exit2 if X12 ≤ X13
```

```
    MOV X0, X21
```

```
// first swap parameter is v
```

```
    MOV X1, X20
```

```
// second swap parameter is j
```

```
    BL swap
```

```
// call swap
```

```
    SUB X20, X20, #1
```

```
// j -= 1
```

```
    B for2tst
```

```
// branch to test of inner loop
```

```
exit2:
```

LDR X12, [X10, X20]
LSL #3

Preserving Registers

■ Preserve saved registers:





```
SUB SP, SP, #40 // make room on stack for 5 regs
STR LR, [SP, #32] // save LR on stack
STR X22, [SP, #24] // save X22 on stack
STR X21, [SP, #16] // save X21 on stack
STR X20, [SP, #8] // save X20 on stack
STR X19, [SP, #0] // save X19 on stack
MOV X21, X0 // copy parameter X0 into X21
MOV X22, X1 // copy parameter X1 into X22
```

■ Restore saved registers:




```
exit1: LDUR X19, [SP, #0] // restore X19 from stack
      LDR X20, [SP, #8] // restore X20 from stack
      LDR X21, [SP, #16] // restore X21 from stack
      LDR X22, [SP, #24] // restore X22 from stack
      LDR X30, [SP, #32] // restore LR from stack
      SUB SP, SP, #40 // restore stack pointer
```

Effect of Compiler Optimization

Compiled with gcc for Pentium 4 under Linux



Effect of Language and Algorithm



Lessons Learnt

- Instruction count and CPI are not good performance indicators in isolation
- Compiler optimizations are sensitive to the algorithm
- Java/JIT compiled code is significantly faster than JVM interpreted
 - Comparable to optimized C in some cases
- Nothing can fix a dumb algorithm!

Arrays vs. Pointers

- Array indexing involves
 - Multiplying index by element size
 - Adding to array base address
- Pointers correspond directly to memory addresses
 - Can avoid indexing complexity

Example: Clearing an Array

```
clear1(int array[], int size) {  
    int i;  
    for (i = 0; i < size; i += 1)  
        array[i] = 0;  
}
```

```
        MOV X9,XZR        // i = 0  
loop1: LSL X10,X9,#3      // X10 = i * 8  
        ADD X11,X0,X10    // X11 = address  
                        // of array[i]  
        STR XZR,[X11,#0]  
                        // array[i] = 0  
        ADD X9,X9,#1      // i = i + 1  
        CMP X9,X1         // compare i to  
                        // size  
        B.LT loop1        // if (i < size)  
                        // go to loop1
```

```
clear2(int *array, int size) {  
    int *p;  
    for (p = &array[0]; p < &array[size];  
        p = p + 1)  
        *p = 0;  
}
```

```
        MOV X9,X0         // p = address of  
                        // array[0]  
        LSL X10,X1,#3      // X10 = size * 8  
        ADD X11,X0,X10    // X11 = address  
                        // of array[size]  
loop2: STR XZR,0[X9,#0]  
                        // Memory[p] = 0  
        ADD X9,X9,#8       // p = p + 8  
        CMP X9,X11         // compare p to <  
                        // &array[size]  
        B.LT loop2        // if (p <  
                        // &array[size])  
                        // go to loop2
```

Comparison of Array vs. Ptr


- Multiply “strength reduced” to shift
- Array version requires shift to be inside loop
 - Part of index calculation for incremented i
 - c.f. incrementing pointer
- Compiler can achieve same effect as manual use of pointers
 - Induction variable elimination
 - Better to make program clearer and safer

ARM & MIPS Similarities

- ARM: the most popular embedded core
- Similar basic set of instructions to MIPS

	ARM	MIPS
Date announced	1985	1985
Instruction size	32 bits	32 bits
Address space	32-bit flat	32-bit flat
Data alignment	Aligned	Aligned
Data addressing modes	9	3
Registers	15 × 32-bit	31 × 32-bit
Input/output	Memory mapped	Memory mapped

Instruction Encoding



The Intel x86 ISA

- Evolution with backward compatibility
 - 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
 - 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
 - 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
 - 80286 (1982): 24-bit addresses, MMU
 - Segmented memory mapping and protection
 - 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments

The Intel x86 ISA

■ Further evolution...

- i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
- Pentium (1993): superscalar, 64-bit datapath
 - Later versions added MMX (Multi-Media eXtension) instructions
 - The infamous FDIV bug
- Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
- Pentium III (1999)
 - Added SSE (Streaming SIMD Extensions) and associated registers
- Pentium 4 (2001)
 - New microarchitecture
 - Added SSE2 instructions

The Intel x86 ISA

- And further...
 - AMD64 (2003): extended architecture to 64 bits
 - EM64T – Extended Memory 64 Technology (2004)
 - AMD64 adopted by Intel (with refinements)
 - Added SSE3 instructions
 - Intel Core (2006)
 - Added SSE4 instructions, virtual machine support
 - AMD64 (announced 2007): SSE5 instructions
 - Intel declined to follow, instead...
 - Advanced Vector Extension (announced 2008)
 - Longer SSE registers, more instructions
- If Intel didn't extend with compatibility, its competitors would!
 - Technical elegance ≠ market success

Basic x86 Registers

Name	31	0	Use
EAX			GPR 0
ECX			GPR 1
EDX			GPR 2
EBX			GPR 3
ESP			GPR 4
EBP			GPR 5
ESI			GPR 6
EDI			GPR 7
	CS		Code segment pointer
	SS		Stack segment pointer (top of stack)
	DS		Data segment pointer 0
	ES		Data segment pointer 1
	FS		Data segment pointer 2
	GS		Data segment pointer 3
EIP			Instruction pointer (PC)
EFLAGS			Condition codes

Basic x86 Addressing Modes


- Two operands per instruction

Source/dest operand	Second source operand
Register	Register
Register	Immediate
Register	Memory
Memory	Register
Memory	Immediate

- Memory addressing modes
 - Address in register
 - $\text{Address} = R_{\text{base}} + \text{displacement}$
 - $\text{Address} = R_{\text{base}} + 2^{\text{scale}} \times R_{\text{index}}$ (scale = 0, 1, 2, or 3)
 - $\text{Address} = R_{\text{base}} + 2^{\text{scale}} \times R_{\text{index}} + \text{displacement}$

x86 Instruction Encoding


a. JE EIP + displacement




b. CALL




c. MOV EBX, [EDI + 45]



d. PUSH ESI



e. ADD EAX, #6765



f. TEST EDX, #42



- Variable length encoding
 - Postfix bytes specify addressing mode
 - Prefix bytes modify operation
 - Operand length, repetition, locking, ...

Implementing IA-32

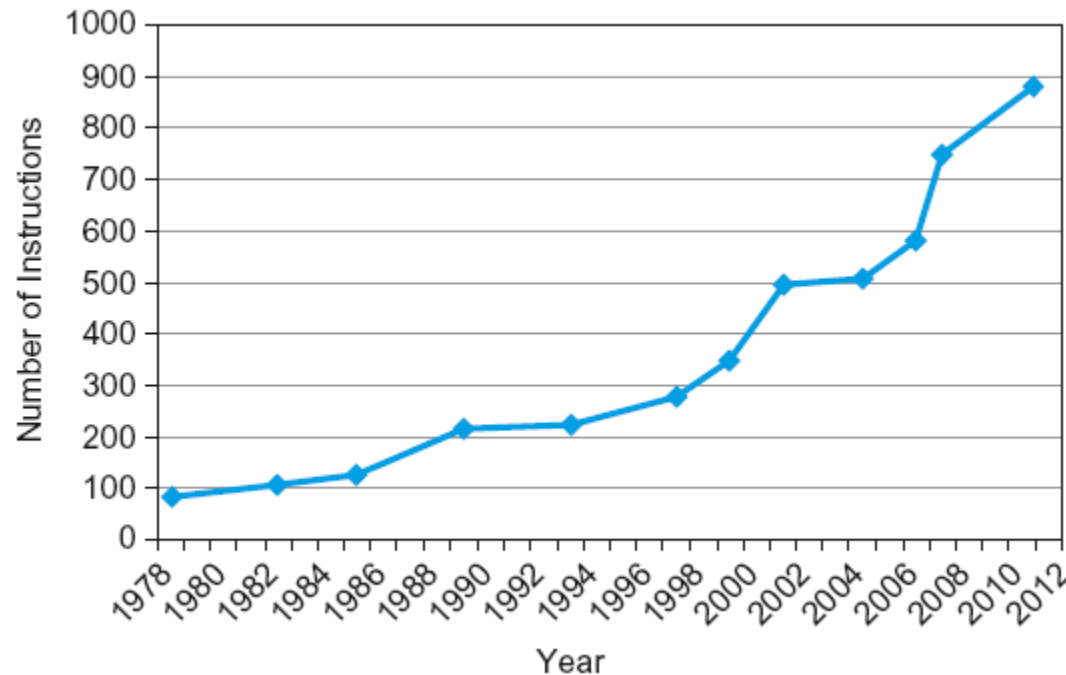
- Complex instruction set makes implementation difficult
 - Hardware translates instructions to simpler microoperations
 - Simple instructions: 1–1
 - Complex instructions: 1–many
 - Microengine similar to RISC
 - Market share makes this economically viable
- Comparable performance to RISC
 - Compilers avoid complex instructions

Fallacies

- Powerful instruction \Rightarrow higher performance
 - Fewer instructions required
 - But complex instructions are hard to implement
 - May slow down all instructions, including simple ones
 - Compilers are good at making fast code from simple instructions
- Use assembly code for high performance
 - But modern compilers are better at dealing with modern processors
 - More lines of code \Rightarrow more errors and less productivity

Fallacies

- Backward compatibility \Rightarrow instruction set doesn't change
 - But they do accrete more instructions



x86 instruction set

Concluding Remarks

- Design principles
 1. Simplicity favors regularity
 2. Smaller is faster
 3. Make the common case fast
 4. Good design demands good compromises
- Layers of software/hardware
 - Compiler, assembler, hardware
- ARMv8: typical of RISC ISAs
 - c.f. x86